

Isolation Forest for Anomaly Detection

Efficient Outlier Detection Using Random Partitioning

Authors: Nicholas Jebeles, Blake Senn, and Isaiah Barlatier

Motivating Questions

Q: Why do we need anomaly detection?

- Fraud detection in financial transactions
- Network intrusion detection
- Manufacturing quality control
- System health monitoring

Q: How can we detect outliers efficiently?

- Traditional methods struggle with large datasets
- Need for unsupervised approaches
- This project demonstrates a solution!

Algorithm Introduction

- **Core Insight:** Anomalies are "few and different"

Key Properties of Anomalies:

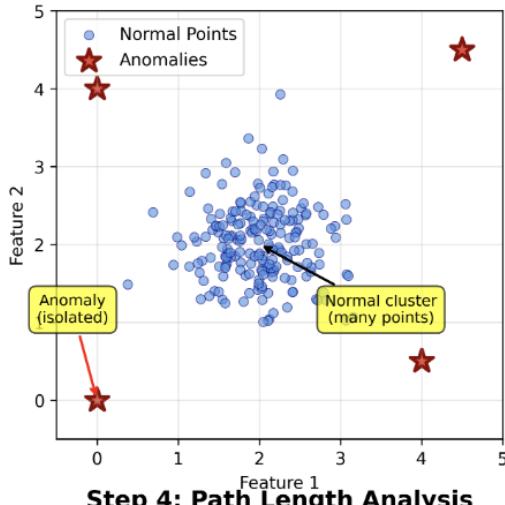
- **Few:** Minority instances in the dataset
- **Different:** Have unusual feature values

The Isolation Forest Approach:

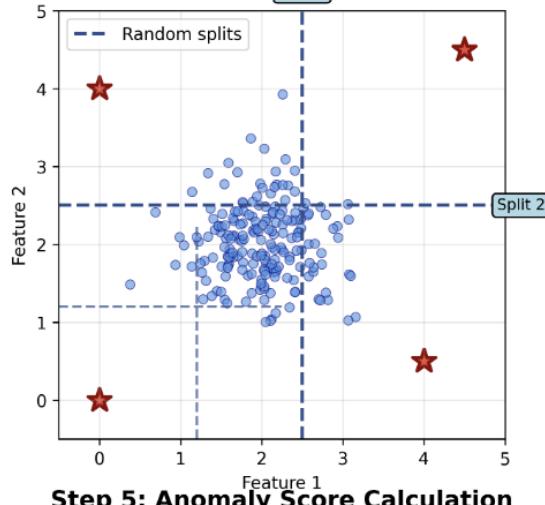
- Uses random partitioning
- Anomalies require fewer splits to isolate
- No distance or density calculations needed

How Isolation Forest Works

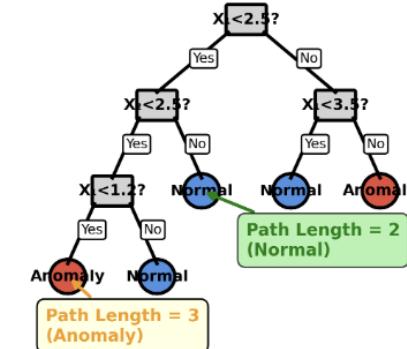
Step 1: Data Distribution



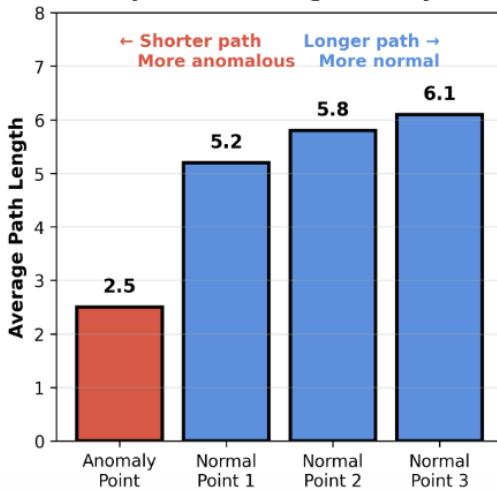
Step 2: Random Partitioning



Step 3: Isolation Tree



Step 4: Path Length Analysis



Step 5: Anomaly Score Calculation

Anomaly Score Formula:

$$S(x, n) = 2^{-\frac{E(h(x))}{c(n)}}$$

Where:

- $E(h(x))$ = Average path length
- $c(n)$ = Normalization constant
- n = Number of samples

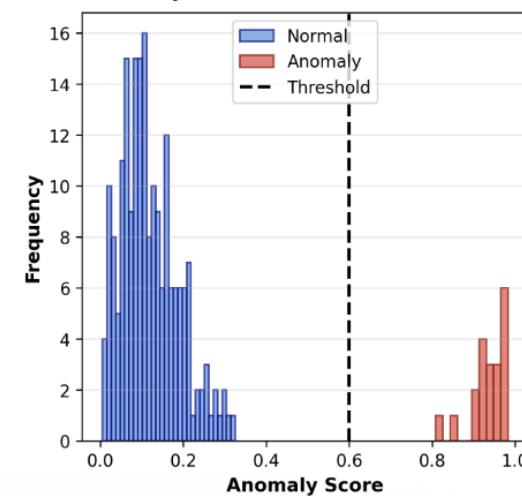
Score Interpretation:

- Score → 1: Anomaly
- Score → 0: Normal
- Score ≈ 0.5: Uncertain

Example Scores:

- Anomaly: $s = 0.85$
- Normal: $s = 0.12$

Step 6: Final Classification



Mathematical Foundation

Path Length Calculation:

$h(x)$ = number of edges from root to isolate point x

$E(h(x))$ = average path length across all trees

Complexity Analysis:

Training: $O(t \times \psi \times \log \psi)$

t = number of trees

ψ = sample size per tree

Prediction: $O(t \times \log \psi)$

Normalization Constant:

$$c(n) = 2H(n-1) - 2(n-1)/n$$

Where $H(i) = \ln(i) + \gamma$ (Euler's constant ≈ 0.5772)

Anomaly Score:

$$s(x, n) = 2^{(-E(h(x)) / c(n))}$$

Interpretation:

- $s(x, n) \rightarrow 1$: Definite anomaly
- $s(x, n) \rightarrow 0$: Definitely normal
- $s(x, n) \approx 0.5$: Ambiguous

Dataset & Preprocessing

Dataset: Credit Card Fraud Detection

- 10,000 transactions
- 20 features
- ~3% fraud rate (highly imbalanced)

Preprocessing Steps:

- Feature scaling (StandardScaler)
- Train/test split (70/30)
- No label information used for training

Model Configuration

Key Parameters:

- **n_estimators:** 100 trees
- **contamination:** 0.03 (3% expected outliers)
- **max_samples:** 'auto' (256 samples per tree)
- **random_state:** 42 (reproducibility)

Why These Values?

- 100 trees: Good balance of accuracy and speed
- Contamination matches actual fraud rate
- Max_samples='auto': Efficient for this data size

Results - Performance Metrics

Performance:

- Precision: 4.9%
- Recall: 3.9%
- F1-Score: 4.3%
- ROC-AUC: 59.2%

Confusion Matrix

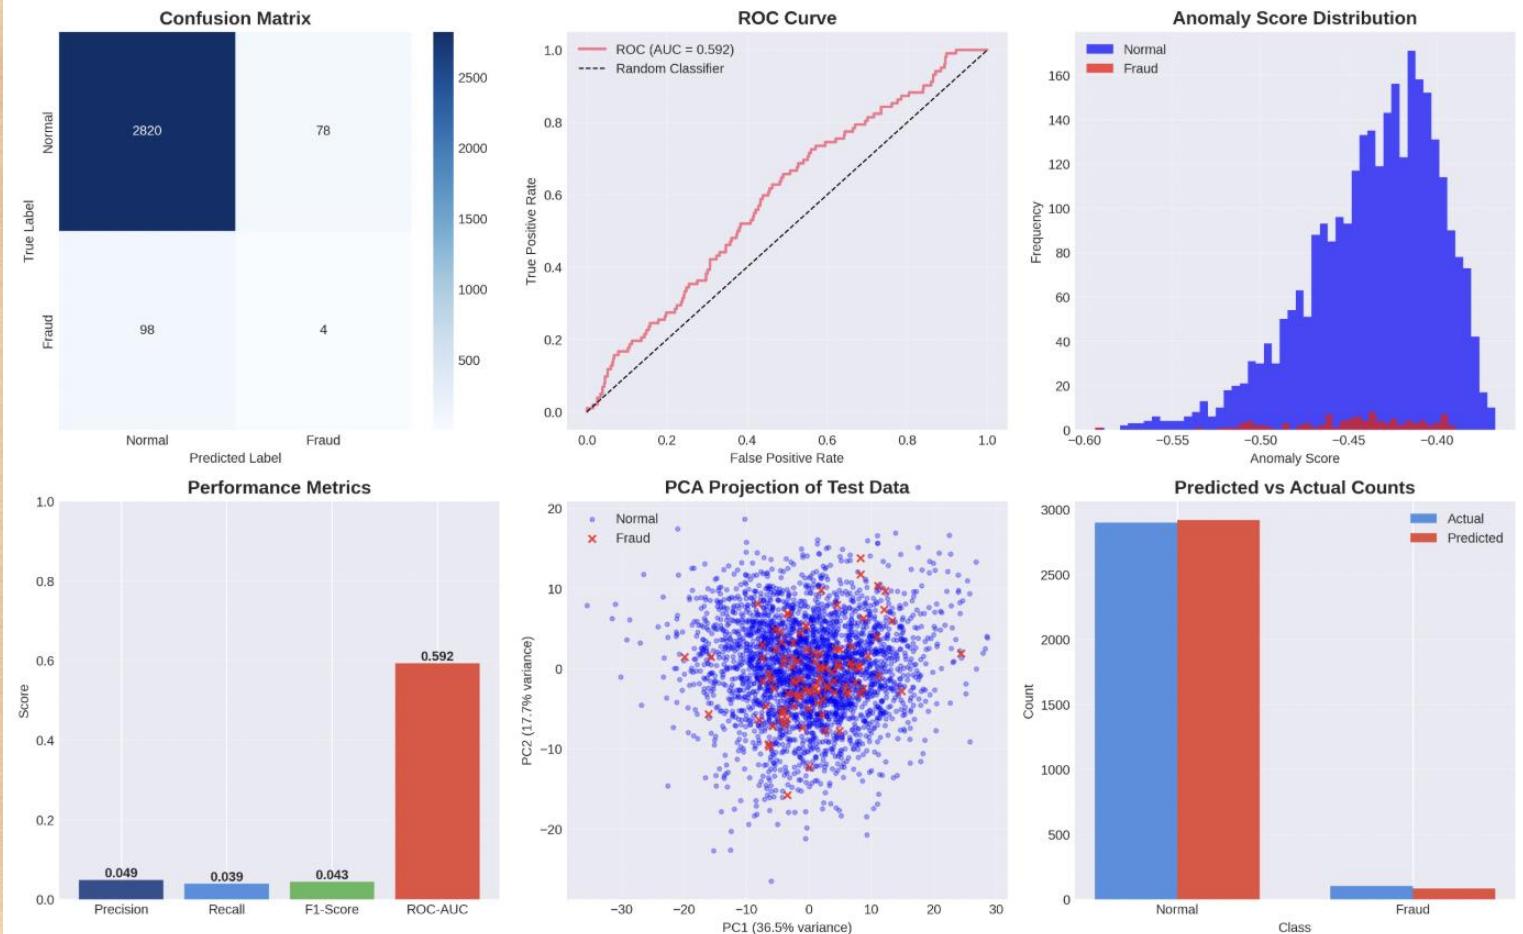
	Predicted Normal	Predicted Fraud
Actual Normal	2,820 (TN)	78 (FP)
Actual Fraud	98 (FN)	4 (TP)

The screenshot shows a Jupyter Notebook interface with a dark theme. On the left is a sidebar with various icons for file operations like opening, saving, and running cells. The main area displays a Python script named `isolation_forest_main.py`. The code is a Mac-compatible version of an Isolation Forest implementation for anomaly detection. It includes imports for numpy, pandas, matplotlib.pyplot, seaborn, sklearn.ensemble, sklearn.model_selection, sklearn.preprocessing, and sklearn.metrics. It also imports PCA from sklearn.decomposition and handles warnings with `warnings.filterwarnings('ignore')`. The script sets a style for better visualizations using `plt.style.use('seaborn-v0_8-darkgrid')` and specifies a color palette with `sns.set_palette("husl")`. The code is annotated with comments explaining its purpose and credits the authors: Blake Senn, Nick Jebelles, and Isaiah Barlatier.

```
File Edit Selection View Go Run ... ← → Search isolation_forest_main.py X C: > Users > blake > Desktop > COMP_5130 > isolation_forest_main.py > ... 1 """ 2 Isolation Forest for Anomaly Detection - Mac Compatible Version 3 ===== 4 This implementation demonstrates anomaly detection using Isolation Forest algorithm 5 for a Data Mining class project. 6 + 7 Author: Blake Senn, Nick Jebelles, Isaiah Barlatier 8 Course: Data Mining 9 Topic: Anomaly Detection with Isolation Forest 10 """ 11 12 import numpy as np 13 import pandas as pd 14 import matplotlib.pyplot as plt 15 import seaborn as sns 16 from sklearn.ensemble import IsolationForest 17 from sklearn.model_selection import train_test_split 18 from sklearn.preprocessing import StandardScaler 19 from sklearn.metrics import ( 20     classification_report, 21     confusion_matrix, 22     roc_auc_score, 23     roc_curve, 24     precision_recall_fscore_support 25 ) 26 from sklearn.decomposition import PCA 27 import warnings 28 warnings.filterwarnings('ignore') 29 30 # Set style for better visualizations 31 plt.style.use('seaborn-v0_8-darkgrid') 32 sns.set_palette("husl")
```

Ln 1, Col 1 Spaces: 4 UTF-8 LF () Python 3.13.7

Results - Visualizations



Code Demonstration

```
# Model initialization
detector = IsolationForestAnomalyDetector(
    n_estimators=100,
    contamination=0.03,
    max_samples='auto'
)
# Training (unsupervised)
detector.fit(X_train)
# Prediction
predictions = detector.predict(X_test)
anomaly_scores = detector.anomaly_scores(X_test)
```

Conclusion & Discussion

- **Key Takeaways:**

- Isolation Forest is efficient and scalable
- No labeled data required
- Strong performance on imbalanced data
- Practical for real-world applications

- **Limitations:**

- May miss local/contextual anomalies
- Requires setting contamination parameter
- Less interpretable than rule-based methods

- **Future Work:**

- Test on multiple datasets
- Compare with other algorithms
- Add feature importance analysis
- Deploy as real-time system